

Import Libraries

```
In [1]: import os
import pickle
import numpy as np
from tqdm.notebook import tqdm

#for preprocessing of images.
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.preprocessing.image import load_img, img_to_array
#for preprocessing of text.
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Model
from tensorflow.keras.utils import to_categorical, plot_model
#Importing layers
from tensorflow.keras.layers import Input, Dense, LSTM, Embedding, Dropout, add
```

Importing Dataset

```
In [2]: BASE_DIR = r'C:\Users\Hp\Desktop\All Documents\Data'
```

Extract Image Features

```
In [3]: # load vgg16 model
model = VGG16()
# Restructure the model
model = Model(inputs=model.inputs, outputs=model.layers[-2].output)
# summarize
print(model.summary())
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312

Total params: 134,260,544
Trainable params: 134,260,544
Non-trainable params: 0

None

```
In [13]: # extract features from image
features = {}
#To get the directory for Images
directory = os.path.join(BASE_DIR, 'Images')

for img_name in tqdm(os.listdir(directory)):
    # load the image from file for getting each image name.
    img_path = directory + '/' + img_name
    image = load_img(img_path, target_size=(224, 224))
```

```

# convert image pixels to numpy array
image = img_to_array(image)
# reshape data for model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
# preprocess image for vgg
image = preprocess_input(image)
# extract features
feature = model.predict(image, verbose=0)
# get image ID
image_id = img_name.split('.')[0]
# store feature
features[image_id] = feature

```

0%| | 0/8091 [00:01<?, ?it/s]

Store and Load Features

```

In [ ]: #Store features in pickle.
pickle.dump(features, open(os.path.join(BASE_DIR, 'features.pkl'), 'wb'))

```

```

In [4]: # load features from pickle
with open(os.path.join(BASE_DIR, 'features.pkl'), 'rb') as f:
    features = pickle.load(f)

```

```

In [5]: model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312

=====
Total params: 134,260,544
Trainable params: 134,260,544
Non-trainable params: 0
=====

Load the Data of Caption

```
In [6]: with open(os.path.join(BASE_DIR, 'captions.txt'), 'r') as f:
        next(f)
        captions_doc = f.read()
```

```
In [8]: # create mapping of image to captions
```

```

# process lines
for line in tqdm(captions_doc.split('\n')):
    # split the line by comma(,)
    tokens = line.split(',')
    if len(line) < 2:
        continue
    image_id, caption = tokens[0], tokens[1:]
    # remove extension from image ID
    image_id = image_id.split('.')[0]
    # convert caption list to string
    caption = " ".join(caption)
    # create list if needed
    if image_id not in mapping:
        mapping[image_id] = []
    # store the caption
    mapping[image_id].append(caption)

```

0%| | 0/40456 [00:00<?, ?it/s]

In [9]: *#To check how many images.*
len(mapping)

Out[9]: 8091

Preprocessing of Captions

In [10]: *#Create function of clean.*

```

def clean(mapping):
    for key, captions in mapping.items():
        for i in range(len(captions)):
            # take one caption at a time
            caption = captions[i]
            # preprocessing steps
            # convert to lowercase
            caption = caption.lower()
            # delete digits, special chars, etc.
            caption = caption.replace('[^A-Za-z]', '')
            # delete additional spaces
            caption = caption.replace('\s+', ' ')
            # add start and end tags to the caption and remove single characters like a
            caption = 'startseq ' + " ".join([word for word in caption.split() if len(word) > 1])
            captions[i] = caption

```

In [11]: *# before preprocess of text*
mapping['1000268201_693b08cb0e']

Out[11]: ['A child in a pink dress is climbing up a set of stairs in an entry way .',
'A girl going into a wooden building .',
'A little girl climbing into a wooden playhouse .',
'A little girl climbing the stairs to her playhouse .',
'A little girl in a pink dress going into a wooden cabin .']

In [12]: *# preprocess the text*
clean(mapping)

In [13]: *# after preprocess of text*
mapping['1000268201_693b08cb0e']

```
Out[13]: ['startseq child in pink dress is climbing up set of stairs in an entry way endseq',
          'startseq girl going into wooden building endseq',
          'startseq little girl climbing into wooden playhouse endseq',
          'startseq little girl climbing the stairs to her playhouse endseq',
          'startseq little girl in pink dress going into wooden cabin endseq']
```

Create Tokenizer and Vocab Size for Text Data

```
In [14]: #Store all captions in single list.
all_captions = []
for key in mapping:
    for caption in mapping[key]:
        all_captions.append(caption)
```

```
In [15]: #To check length of captions.
len(all_captions)
```

```
Out[15]: 40455
```

```
In [16]: all_captions[:10]
```

```
Out[16]: ['startseq child in pink dress is climbing up set of stairs in an entry way endseq',
          'startseq girl going into wooden building endseq',
          'startseq little girl climbing into wooden playhouse endseq',
          'startseq little girl climbing the stairs to her playhouse endseq',
          'startseq little girl in pink dress going into wooden cabin endseq',
          'startseq black dog and spotted dog are fighting endseq',
          'startseq black dog and tri-colored dog playing with each other on the road endseq',
          'startseq black dog and white dog with brown spots are staring at each other in the street endseq',
          'startseq two dogs of different breeds looking at each other on the road endseq',
          'startseq two dogs on pavement moving toward each other endseq']
```

```
In [17]: # tokenize the text
#Initialize the tokenizer
tokenizer = Tokenizer()
tokenizer.fit_on_texts(all_captions)
vocab_size = len(tokenizer.word_index) + 1
```

```
In [18]: vocab_size
```

```
Out[18]: 8485
```

```
In [19]: # get maximum length of the caption available
max_length = max(len(caption.split()) for caption in all_captions)
max_length
```

```
Out[19]: 35
```

Train Test Split

```
In [20]: #For getting image ids.
image_ids = list(mapping.keys())
split = int(len(image_ids) * 0.90)
split
train = image_ids[:split]
test = image_ids[split:]
```

Data Generator

```

In [21]: # create data generator to get data in batch (avoids session crash)
def data_generator(data_keys, mapping, features, tokenizer, max_length, vocab_size, batch_size):
    # loop over images
    X1, X2, y = list(), list(), list()
    #for determining whether we reach the batch size or not.
    n = 0
    while 1:
        for key in data_keys:
            n += 1
            captions = mapping[key]
            # process each caption
            for caption in captions:
                # encode the sequence, this will assign each word an index.
                seq = tokenizer.texts_to_sequences([caption])[0]
                # split the sequence into X (input), y (output) pairs.
                for i in range(1, len(seq)):
                    # split into input and output pairs
                    in_seq, out_seq = seq[:i], seq[i]
                    # pad input sequence
                    in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
                    # encode output sequence
                    out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]

                # store the sequences
                X1.append(features[key][0])
                X2.append(in_seq)
                y.append(out_seq)
            if n == batch_size:
                X1, X2, y = np.array(X1), np.array(X2), np.array(y)
                #return the collected samples to generator.
                yield [X1, X2], y
                X1, X2, y = list(), list(), list()
                n = 0

```

```

In [89]: # encoder model
# image feature layers
!pip install pydot

```

Collecting pydot

Downloading pydot-1.4.2-py2.py3-none-any.whl (21 kB)

Requirement already satisfied: pyparsing>=2.1.4 in c:\users\hp\anaconda3\lib\site-packages (from pydot) (3.0.4)

Installing collected packages: pydot

Successfully installed pydot-1.4.2

('You must install pydot (`pip install pydot`) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) ', 'for plot_model/model_to_dot to work.')

Model Creation

```

In [22]: import matplotlib.pyplot as plt
import graphviz
import pydot

#Encoder Model
#Image feature layers
#4096 is the shape of features.
inputs1 = Input(shape=(4096))
fe1 = Dropout(0.4)(inputs1)
fe2 = Dense(256, activation='relu')(fe1)
# sequence feature layers
inputs2 = Input(shape=(max_length))
e2 = Embedding(vocab_size, 256, mask_zero=True)(inputs2)

```

```

#for avoiding overfitting we use 0.4 or 0.5.
se2 = Dropout(0.4)(se1)
se3 = LSTM(256)(se2)

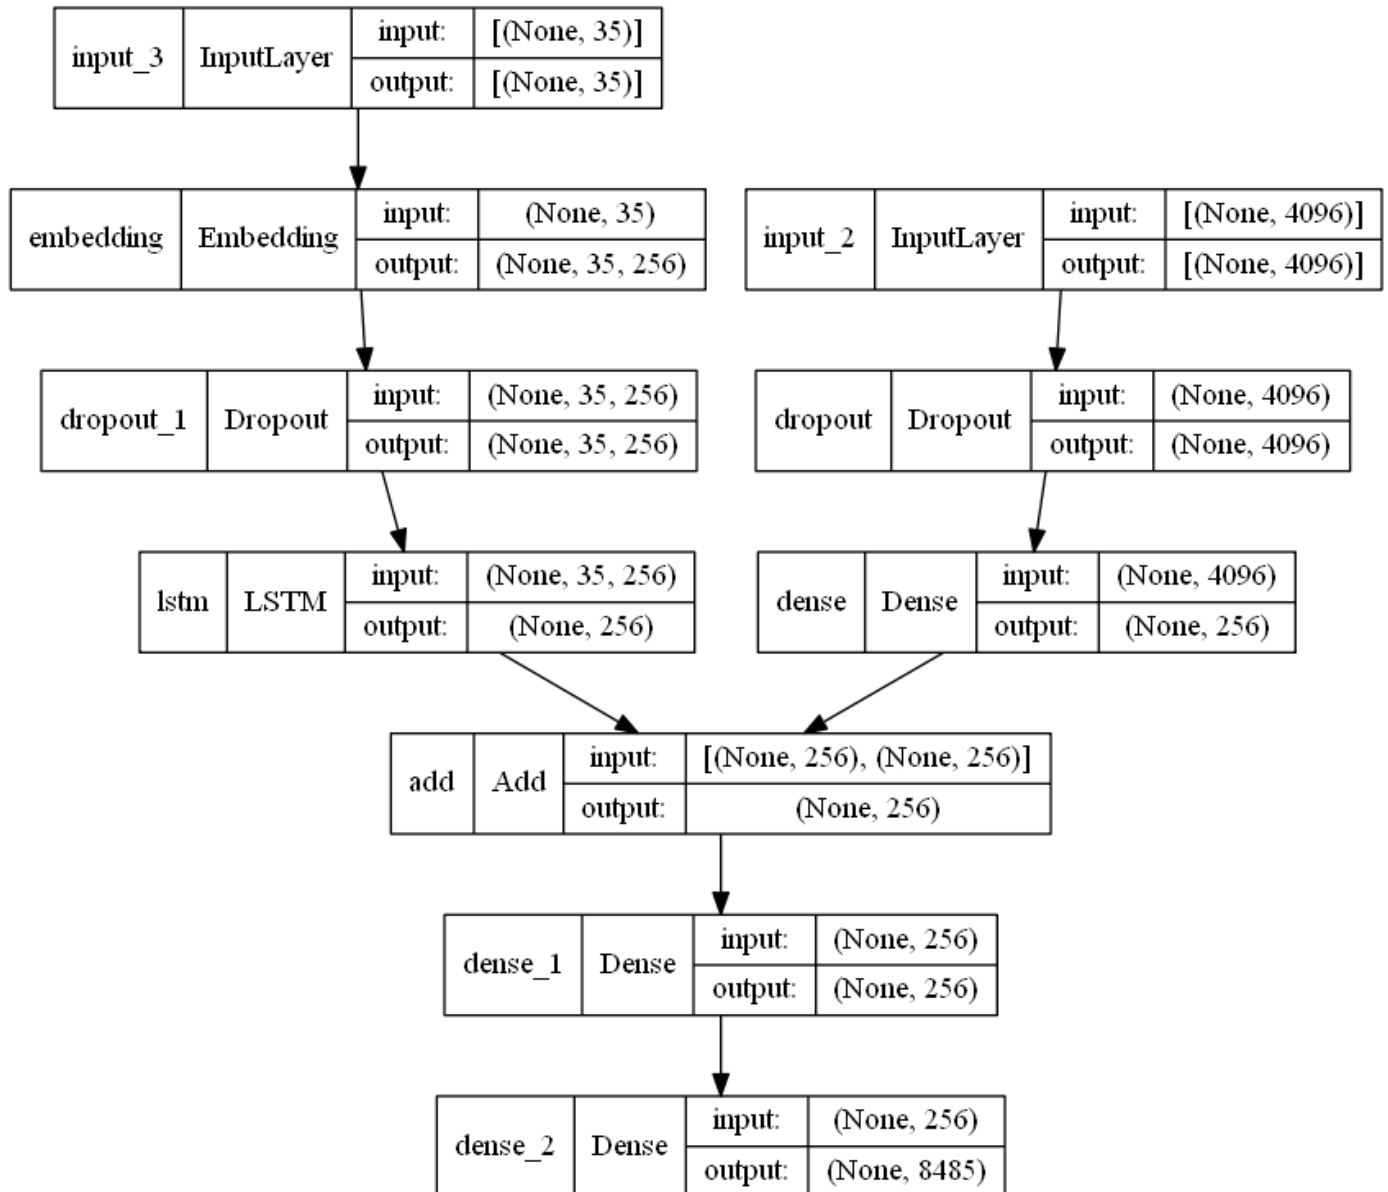
# Decoder Model
#concatinating the image and text features.
decoder1 = add([fe2, se3])
decoder2 = Dense(256, activation='relu')(decoder1)
outputs = Dense(vocab_size, activation='softmax')(decoder2)

model = Model(inputs=[inputs1, inputs2], outputs=outputs)
model.compile(loss='categorical_crossentropy', optimizer='adam')

# plot the model
plot_model(model, show_shapes=True)

```

Out[22]:



In [23]: model.summary()

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 35)]	0	[]
input_2 (InputLayer)	[(None, 4096)]	0	[]
embedding (Embedding)	(None, 35, 256)	2172160	['input_3[0][0]']
dropout (Dropout)	(None, 4096)	0	['input_2[0][0]']
dropout_1 (Dropout)	(None, 35, 256)	0	['embedding[0][0]']
dense (Dense)	(None, 256)	1048832	['dropout_1[0][0]']
lstm (LSTM)	(None, 256)	525312	['dropout_1[0][0]']
add (Add)	(None, 256)	0	['dense[0][0]', 'lstm[0][0]']
dense_1 (Dense)	(None, 256)	65792	['add[0][0]']
dense_2 (Dense)	(None, 8485)	2180645	['dense_1[0][0]']
Total params: 5,992,741			
Trainable params: 5,992,741			
Non-trainable params: 0			

Train the Model

```
In [24]: # train the model
epochs = 20
batch_size = 32
steps = len(train) // batch_size

for i in range(epochs):
    # create data generator
    generator = data_generator(train, mapping, features, tokenizer, max_length, vocab_si
    # fit for one epoch
    model.fit(generator, epochs=1, steps_per_epoch=steps, verbose=1)
```

```

227/227 [=====] - 1816s 8s/step - loss: 5.2080
227/227 [=====] - 1792s 8s/step - loss: 4.0111
227/227 [=====] - 1776s 8s/step - loss: 3.5766
227/227 [=====] - 1791s 8s/step - loss: 3.3080
227/227 [=====] - 1798s 8s/step - loss: 3.1108
227/227 [=====] - 1788s 8s/step - loss: 2.9623
227/227 [=====] - 1784s 8s/step - loss: 2.8463
227/227 [=====] - 1794s 8s/step - loss: 2.7546
227/227 [=====] - 1803s 8s/step - loss: 2.6744
227/227 [=====] - 1791s 8s/step - loss: 2.6035
227/227 [=====] - 1803s 8s/step - loss: 2.5378
227/227 [=====] - 1795s 8s/step - loss: 2.4787
227/227 [=====] - 1787s 8s/step - loss: 2.4250
227/227 [=====] - 1792s 8s/step - loss: 2.3776
227/227 [=====] - 1788s 8s/step - loss: 2.3400
227/227 [=====] - 1793s 8s/step - loss: 2.3008
227/227 [=====] - 2785s 12s/step - loss: 2.2662
227/227 [=====] - 1779s 8s/step - loss: 2.2419
227/227 [=====] - 1785s 8s/step - loss: 2.2078
227/227 [=====] - 1786s 8s/step - loss: 2.1818

```

Save the Model

```

In [70]: import os.path
         if os.path.isfile('BASE_DIR/save_model.h5') is False:
             model.save('BASE_DIR/save_model.h5')

```

Load the Model

```

In [24]: from tensorflow.keras.models import load_model
         new_model = load_model('BASE_DIR/save_model.h5')

```

```

In [25]: new_model.summary()

```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 35)]	0	[]
input_2 (InputLayer)	[(None, 4096)]	0	[]
embedding (Embedding)	(None, 35, 256)	2172160	['input_3[0][0]']
dropout (Dropout)	(None, 4096)	0	['input_2[0][0]']
dropout_1 (Dropout)	(None, 35, 256)	0	['embedding[0][0]']
dense (Dense)	(None, 256)	1048832	['dropout[0][0]']
lstm (LSTM)	(None, 256)	525312	['dropout_1[0][0]']
add (Add)	(None, 256)	0	['dense[0][0]', 'lstm[0][0]']
dense_1 (Dense)	(None, 256)	65792	['add[0][0]']
dense_2 (Dense)	(None, 8485)	2180645	['dense_1[0][0]']
Total params: 5,992,741			
Trainable params: 5,992,741			
Non-trainable params: 0			

Generate Caption for Images

```
In [26]: #Convert index to words.
def idx_to_word(integer, tokenizer):
    for word, index in tokenizer.word_index.items():
        if index == integer:
            return word
    return None
```

```
In [27]: # generate caption for an image
def predict_caption(new_model, image, tokenizer, max_length):
    # add start tag for generation process
    in_text = 'startseq'
    # iterate over the max length of sequence
    for i in range(max_length):
        # encode input sequence
        sequence = tokenizer.texts_to_sequences([in_text])[0]
        # pad the sequence
        sequence = pad_sequences([sequence], max_length)
        # predict next word
        yhat = new_model.predict([image, sequence], verbose=0)
        # get index with high probability
        yhat = np.argmax(yhat)
        # convert index to word
        word = idx_to_word(yhat, tokenizer)
        # stop if word not found
        if word is None:
```

```

        break
        # append word as input for generating next word
        in_text += " " + word
        # stop if we reach end tag
        if word == 'endseq':
            break

    return in_text

```

Validation with Test data and Bleu Score

```

In [28]: from nltk.translate.bleu_score import corpus_bleu
        # validate with test data
        actual, predicted = list(), list()

        for key in tqdm(test):
            # get actual caption
            captions = mapping[key]
            # predict the caption for image
            y_pred = predict_caption(new_model, features[key], tokenizer, max_length)
            # split into words
            actual_captions = [caption.split() for caption in captions]
            y_pred = y_pred.split()
            # append to the list
            actual.append(actual_captions)
            predicted.append(y_pred)

        # calculate BLEU score
        print("BLEU-1: %f" % corpus_bleu(actual, predicted, weights=(1.0, 0, 0, 0)))
        print("BLEU-2: %f" % corpus_bleu(actual, predicted, weights=(0.5, 0.5, 0, 0)))

0%|          | 0/810 [00:00<?, ?it/s]
BLEU-1: 0.530205
BLEU-2: 0.308311

```

Visualize the Results

```

In [29]: #Image for loading the image
        from PIL import Image
        import matplotlib.pyplot as plt
        def generate_caption(image_name):
            # load the image
            # image_name = "1001773457_577c3a7d70.jpg"
            image_id = image_name.split('.')[0]
            img_path = os.path.join(BASE_DIR, "Images", image_name)
            image = Image.open(img_path)
            captions = mapping[image_id]
            print('-----Actual-----')
            for caption in captions:
                print(caption)
            # predict the caption
            y_pred = predict_caption(new_model, features[image_id], tokenizer, max_length)
            print('-----Predicted-----')
            print(y_pred)
            plt.imshow(image)

```

```

In [30]: generate_caption("1000268201_693b08cb0e.jpg")

```

```

-----Actual-----
startseq child in pink dress is climbing up set of stairs in an entry way endseq
startseq girl going into wooden building endseq
startseq little girl climbing into wooden playhouse endseq
startseq little girl climbing the stairs to her playhouse endseq
startseq little girl in pink dress going into wooden cabin endseq
-----Predicted-----
startseq little girl in red dress is climbing the stairs endseq

```




```
In [31]: generate_caption("1001773457_577c3a7d70.jpg")
```

```

-----Actual-----
startseq black dog and spotted dog are fighting endseq
startseq black dog and tri-colored dog playing with each other on the road endseq
startseq black dog and white dog with brown spots are staring at each other in the street endseq
startseq two dogs of different breeds looking at each other on the road endseq
startseq two dogs on pavement moving toward each other endseq
-----Predicted-----
startseq two dogs playing on the street endseq

```



```
In [32]: generate_caption("1007320043_627395c3d8.jpg")
```

```

-----Actual-----
startseq child playing on rope net endseq
startseq little girl climbing on red roping endseq
startseq little girl in pink climbs rope bridge at the park endseq
startseq small child grips onto the red ropes at the playground endseq
startseq the small child climbs on red ropes on playground endseq
-----Predicted-----
startseq little girl grips the equipment on the playground endseq

```



In [33]: `generate_caption("1009434119_febe49276a.jpg")`

```
-----Actual-----
startseq black and white dog is running in grassy garden surrounded by white fence endseq
q
startseq black and white dog is running through the grass endseq
startseq boston terrier is running in the grass endseq
startseq boston terrier is running on lush green grass in front of white fence endseq
startseq dog runs on the green grass near wooden fence endseq
-----Predicted-----
startseq dog jumps over an obstacle endseq
```

